

Strategies for Algorithm Design

Greedy Algorithms

This week

- Greedy algorithms!
- Builds on our ideas from dynamic programming

Example

Activity selection

You can only do one activity at a time, and you want to maximize the number of activities that you do.

What to choose?

CS
Section

Sleep

Seminar

Program
team

5 Class

Underwater basket
weaving class

CS study
group

Swimming
lessons

Combinatorics
Seminar

Theory Lunch

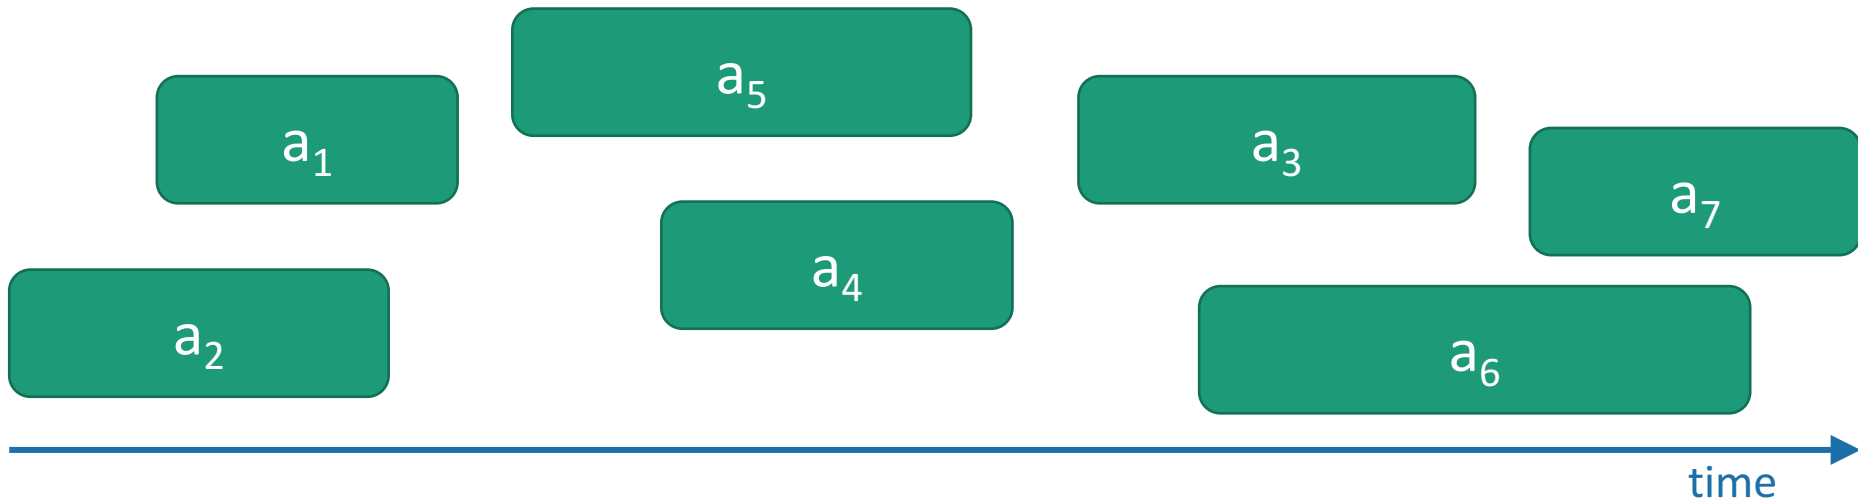
Social activity

time

Activity selection

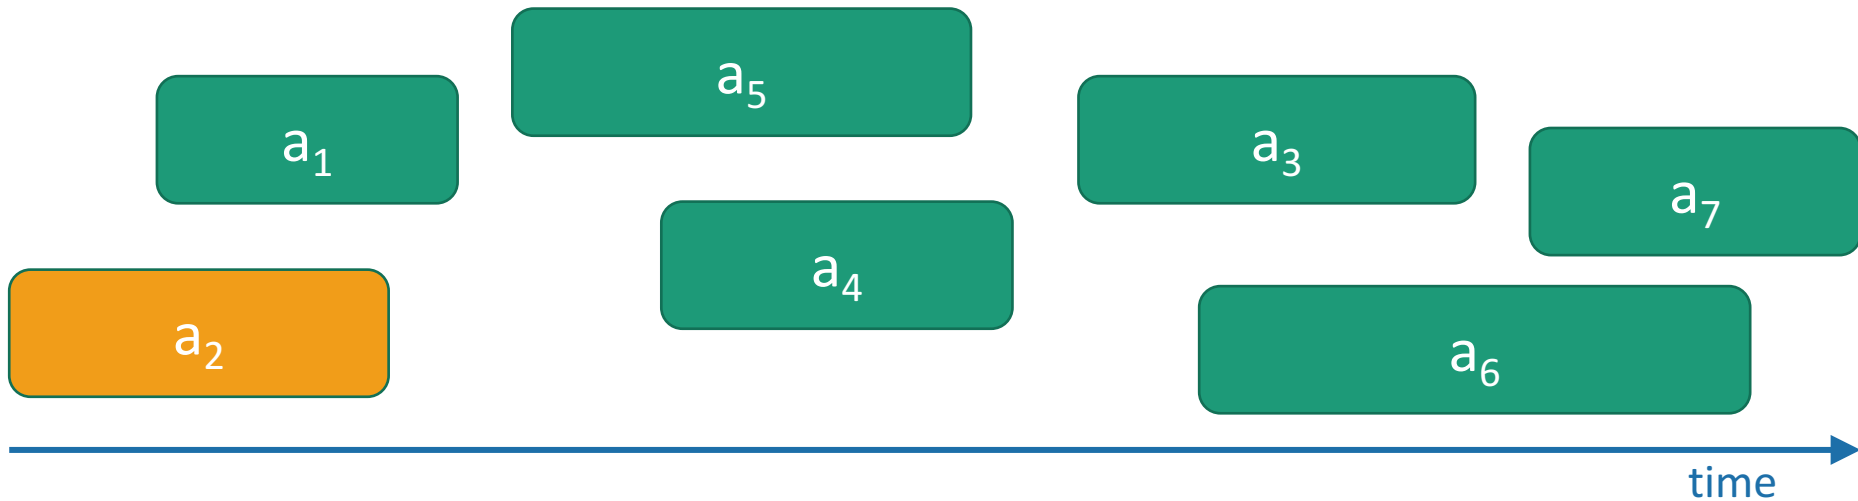
- Input:
 - Activities a_1, a_2, \dots, a_n
 - Start times s_1, s_2, \dots, s_n
 - Finish times f_1, f_2, \dots, f_n
- Output:
 - How many activities can you do today?

Greedy Algorithm



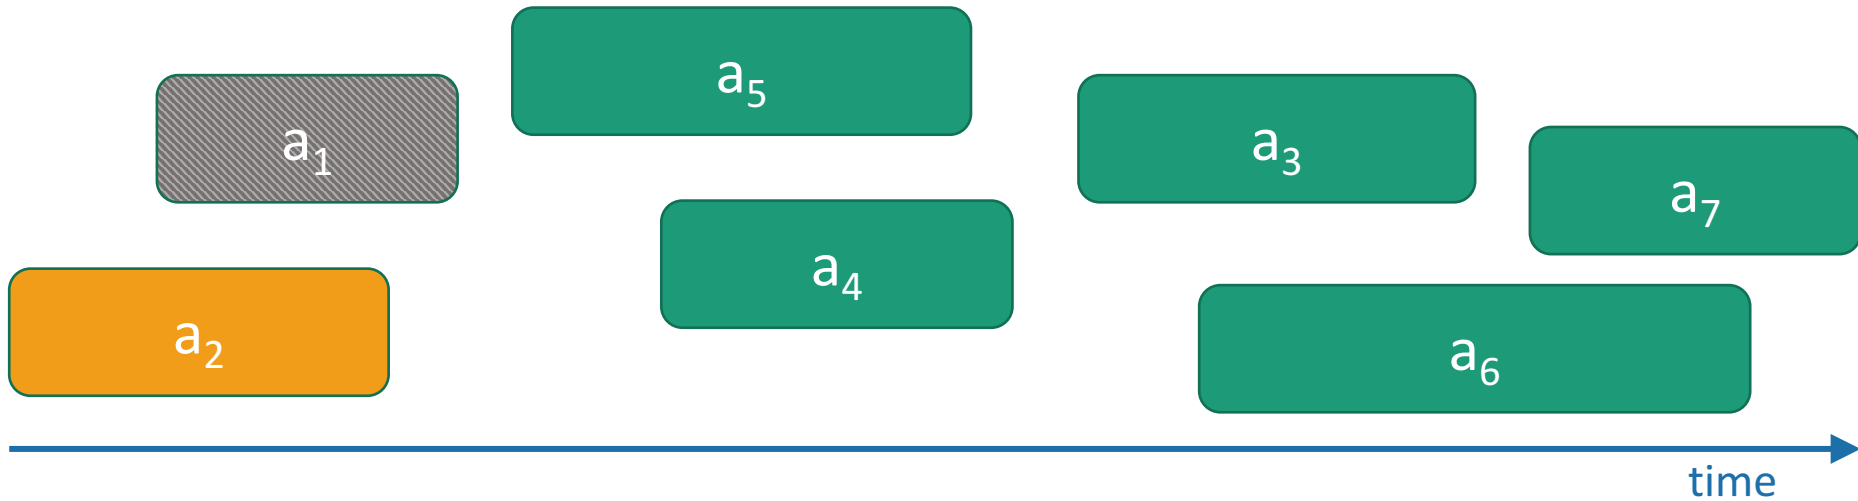
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm




- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm




- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



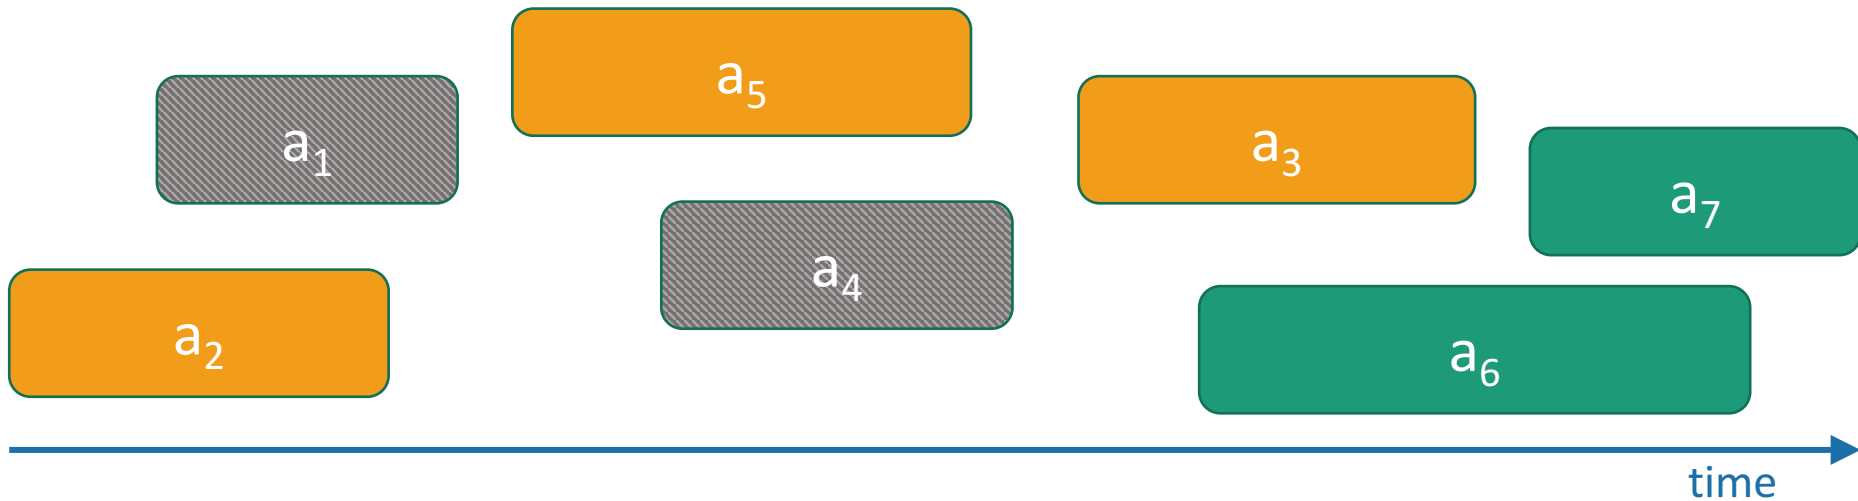
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



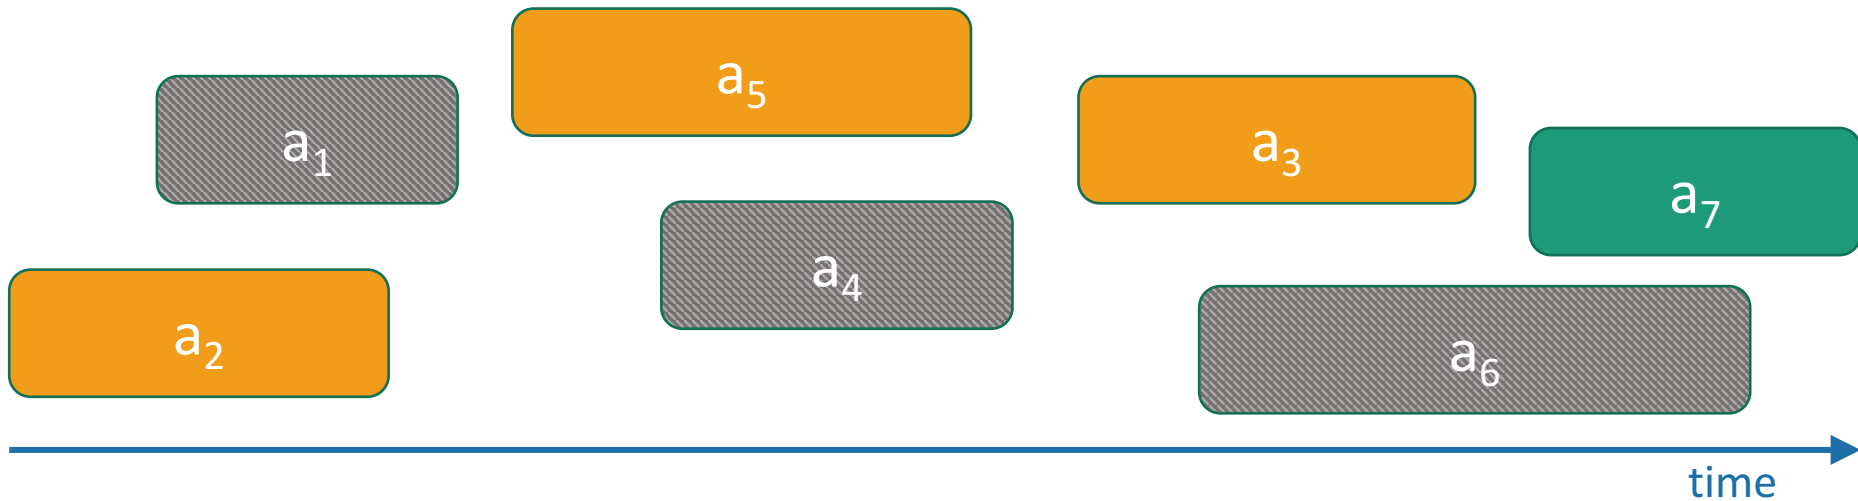
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm




- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- Repeat.

Greedy Algorithm



- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- Repeat.

That seems like a reasonable thing to do...

- Running time:
 - $O(n)$ if the activities are already sorted by finish time.
 - Otherwise $O(n\log(n))$ if you have to sort them first.
- Does it work?
 - We'll see soon.

This is an example of a greedy algorithm

- At each step in the algorithm, make a choice.
 - Hey, I can increase my activity set by one,
 - And leave lots of room for future choices,
 - Let's do that and hope for the best!!!
- **Hope** that at the end of the day, this results in a globally optimal solution.

Three questions


1. Does this greedy algorithm for activity selection work?
2. In general, when are greedy algorithms a good idea?
3. The “greedy” approach is often the first you’d think of...

Answers

1. Does this greedy algorithm for activity selection work?
 - Yes.
2. In general, when are greedy algorithms a good idea?
 - When they exhibit especially nice optimal substructure.
3. The “greedy” approach is often the first you’d think of...

DP view of activity selection

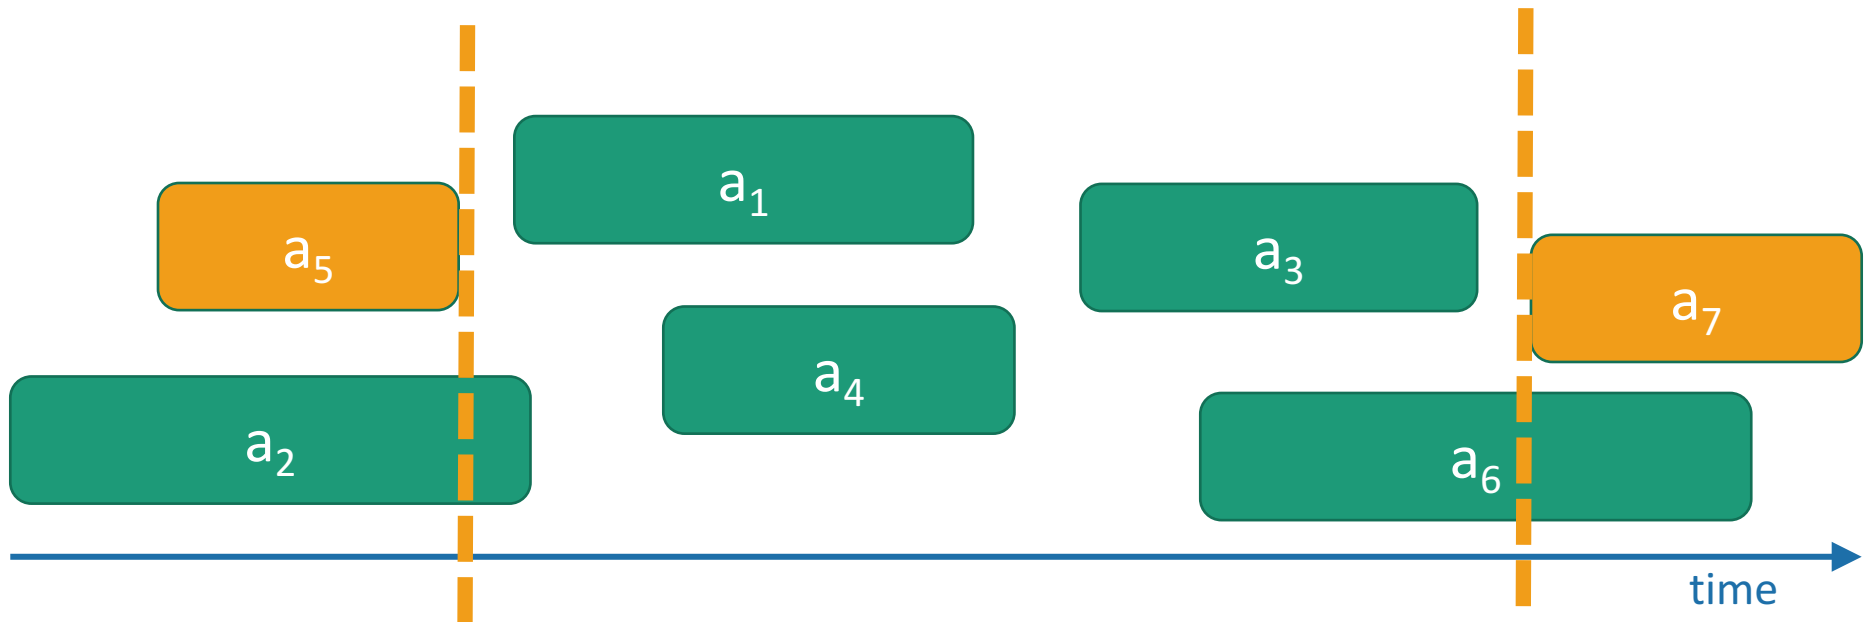
Recipe for applying Dynamic Programming

- **Step 1:** Identify optimal substructure. 
- **Step 2:** Find a recursive formulation for the value of the optimal solution.
- **Step 3:** Use dynamic programming to find the value of the optimal solution.
- **Step 4:** If needed, keep track of some additional info so that the algorithm from Step 3 can find the actual solution.
- **Step 5:** If needed, code this up like a reasonable person.

Optimal substructure

- Subproblems:


number of activities you can squeeze in after Activity i finishes and before Activity j starts



Optimal substructure

- Subproblems:

number of activities you can squeeze in after Activity i finishes and before Activity j starts




Now let's define an optimal solution, i.e., a maximal set of mutually compatible activities between i, j .

Details for optimal substructure property

Then the candidate activities to consider are those that start after a_i and end before a_j :

Now let's define A_{ij} to be an optimal solution, i.e., a maximal set of mutually compatible activities in S_{ij} . What is the structure of this solution?

At some point we will need to make a choice to include some activity a_k , two sets of compatible candidates after a_k is taken out:



S_{ik} : activities that start after a_i finishes, and finish before a_k starts
 S_{kj} : activities that start after a_k finishes, and finish before a_j starts

Note that S_{ij} may be a proper superset of $S_{ik} \cup \{a_k\} \cup S_{kj}$, as activities incompatible with a_k are excluded.)

Using the same notation as above, define the optimal solutions to these subproblems to be:

$A_{ik} = A_{ij} \cap S_{ik}$: the optimal solution to S_{ik}

$A_{kj} = A_{ij} \cap S_{kj}$: the optimal solution to S_{kj}

So the structure of an optimal solution A_{ij} is:

$$A_{ij} = A_{ik} \cup \{a_k\} \cup A_{kj}$$

Recipe for applying Dynamic Programming

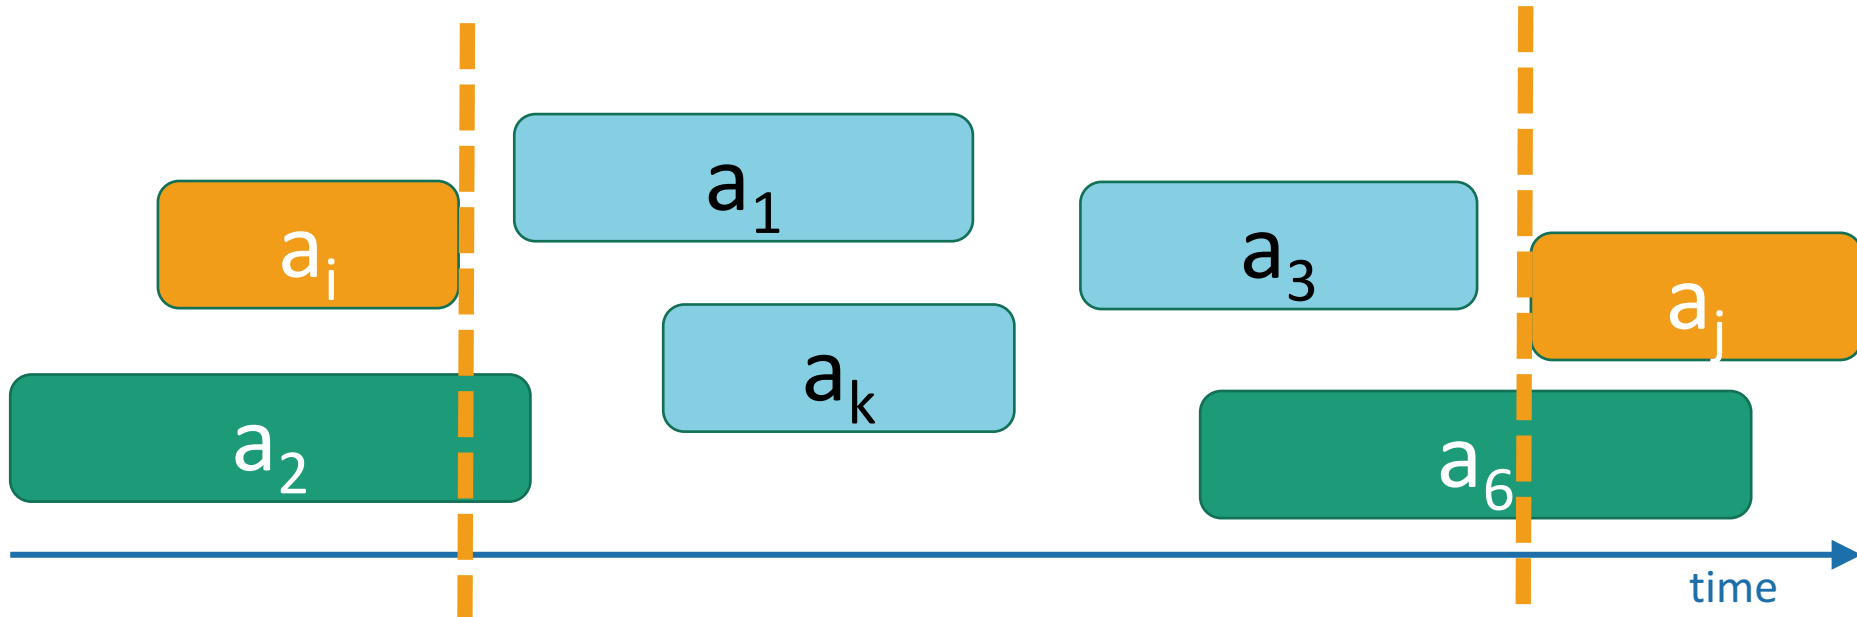
- **Step 1:** Identify optimal substructure.
- **Step 2:** Find a recursive formulation for the value of the optimal solution.
- **Step 3:** Use dynamic programming to find the value of the optimal solution.
- **Step 4:** If needed, keep track of some additional info so that the algorithm from Step 3 can find the actual solution.
- **Step 5:** If needed, code this up like a reasonable person.



This satisfies a nice recursive relationship

$A[i,j]$ = number of activities you can squeeze in after Activity i finishes and before Activity j starts

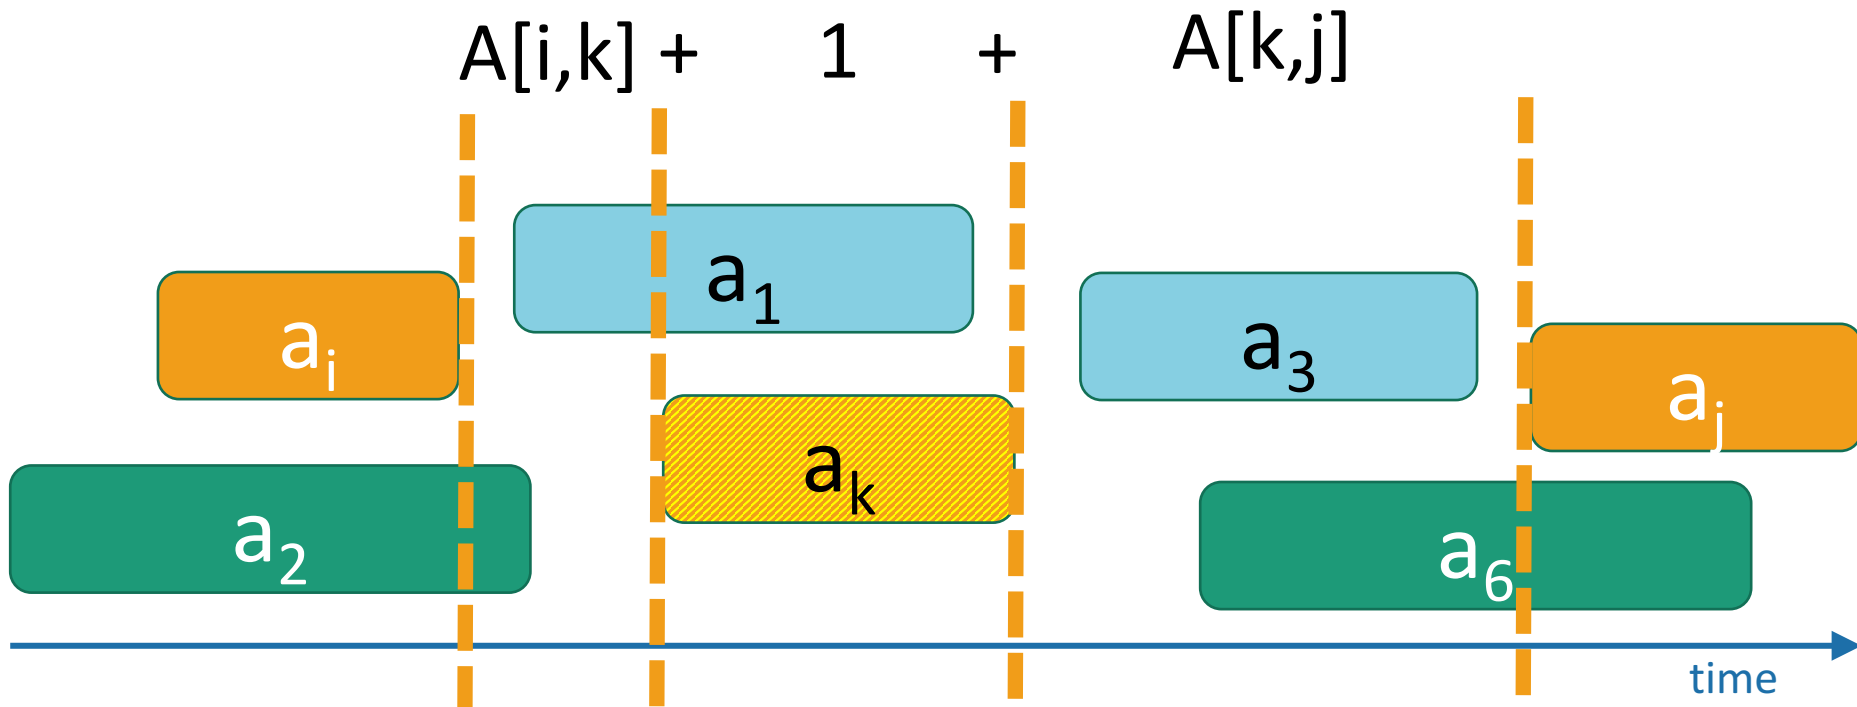
- $A[i,j] = \max_k \{ A[i,k] + 1 + A[k,j] \}$
 - The maximum is over all k so that Activity k fits in between Activities i and j .



This satisfies a nice recursive relationship

$A[i,j]$ = number of activities you can squeeze in after Activity i finishes and before Activity j starts


- $A[i,j] = \max_k \{ A[i,k] + 1 + A[k,j] \}$
 - The maximum is over all k so that Activity k fits in between Activities i and j .



We could turn this into a DP algorithm

- .Would take time something like $O(n^3)$
 - Fill out an n -by- n table.
 - For each entry search over maybe n possibilities for k .
- But this would be wasteful!
 - we just saw an algorithm that takes time $O(n \log(n))$, if it's correct...

Try it!
It builds character!



$A[i,j]$ = number of activities you can squeeze in after Activity i finishes and before Activity j starts

The thing that's wasteful

- Actually, we **should know in advance** what subproblem to look at.
- Lemma:
 - Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
 - Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .

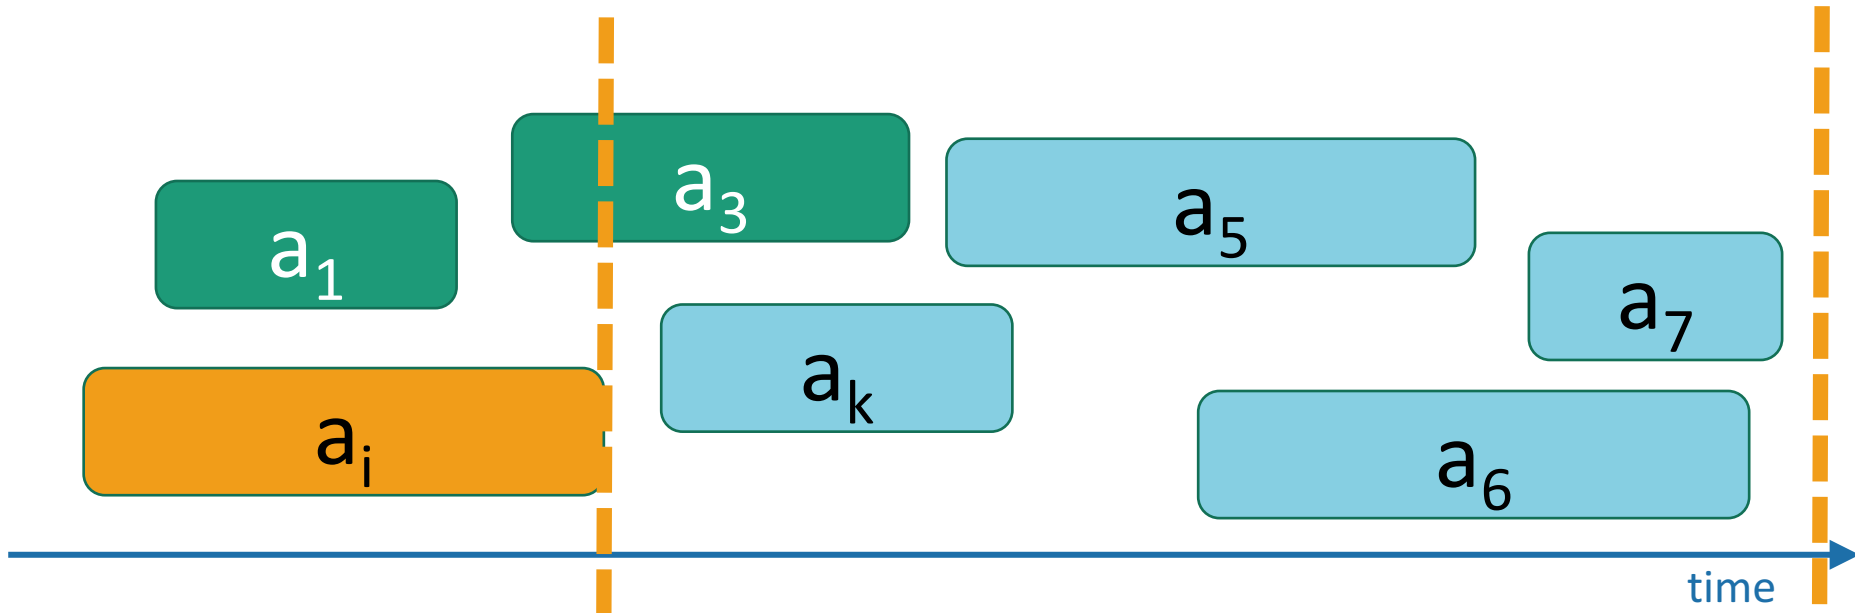
This is abusing notation...technically $A[i..n+1]$ is a value, not a problem.

Let's add an additional activity a_{n+1} that starts "tomorrow".

$A[i,j]$ = number of activities you can squeeze in after Activity i finishes and before Activity j starts

Lemma

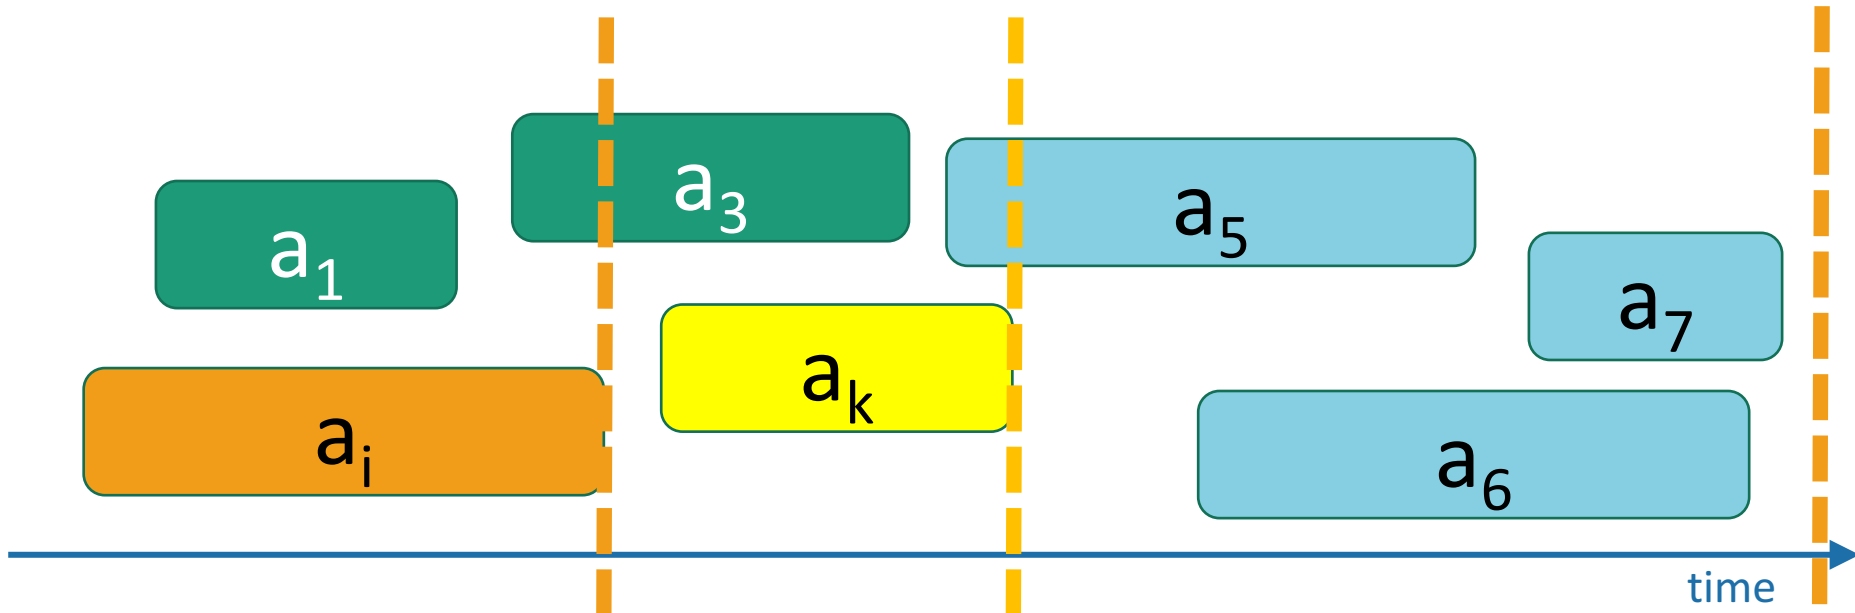
- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .



$A[i,j]$ = number of activities you can squeeze in after Activity i finishes and before Activity j starts

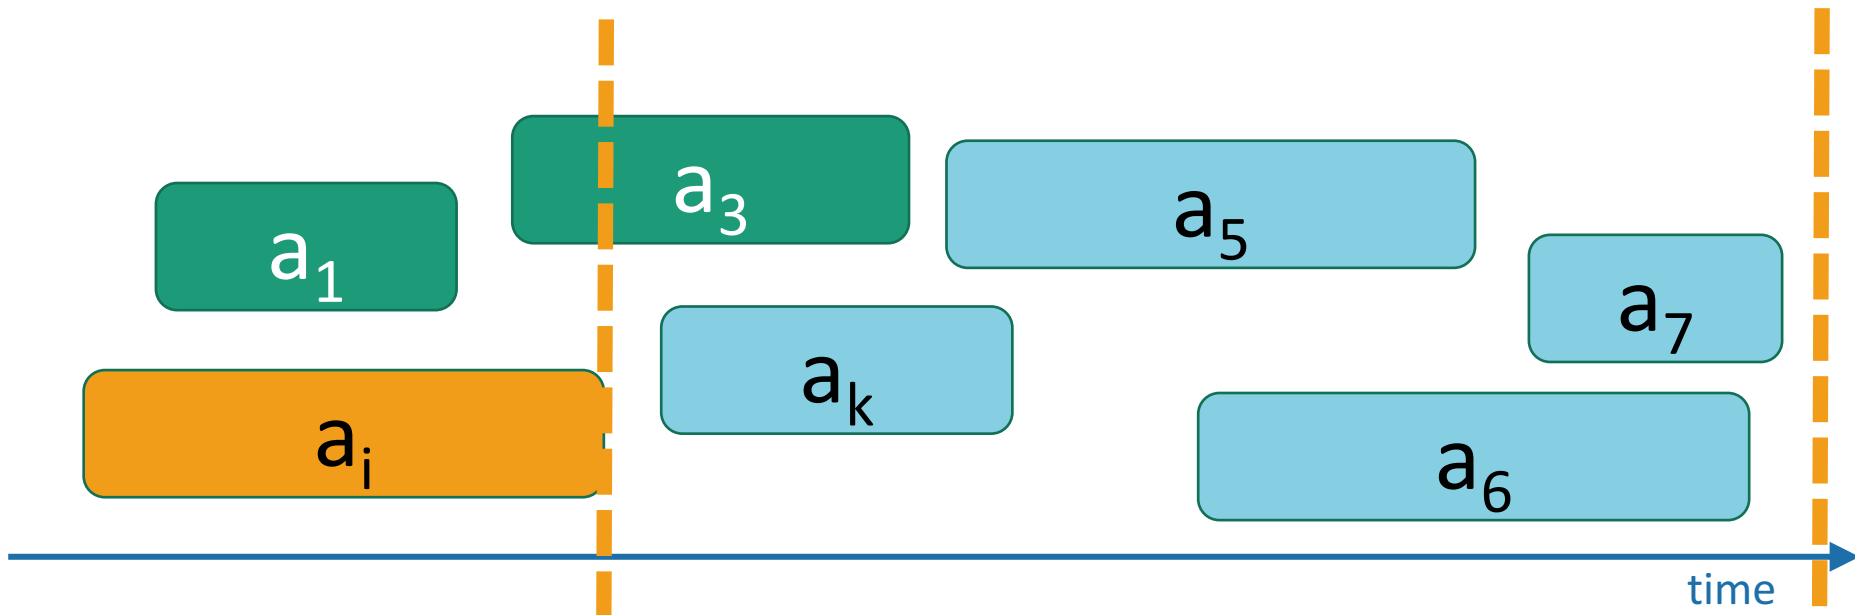
Lemma

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .



Proof


- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .



Proof

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .

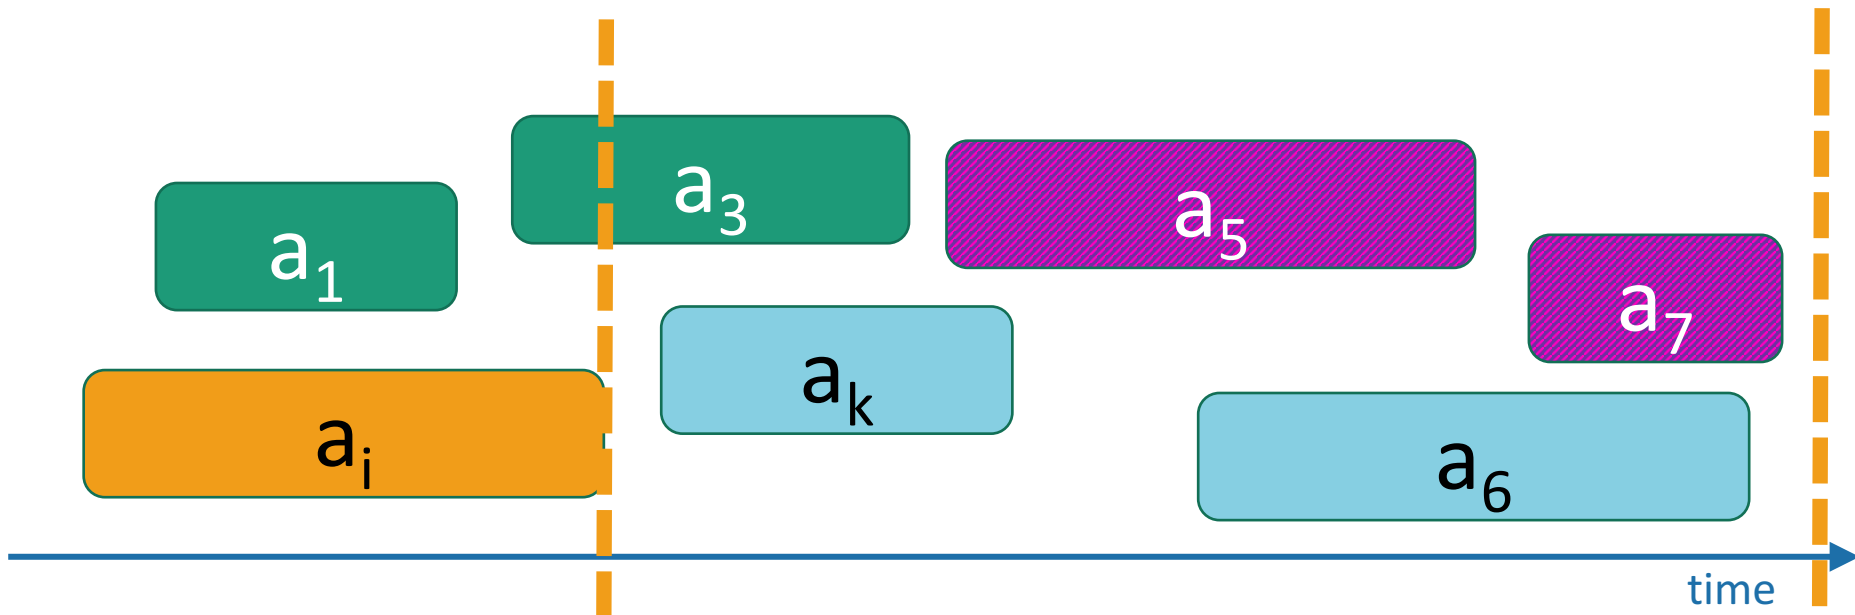
- Suppose that **this** is an optimal solution to **$A[i..n+1]$**
 - Doesn't involve a_k



Proof

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .

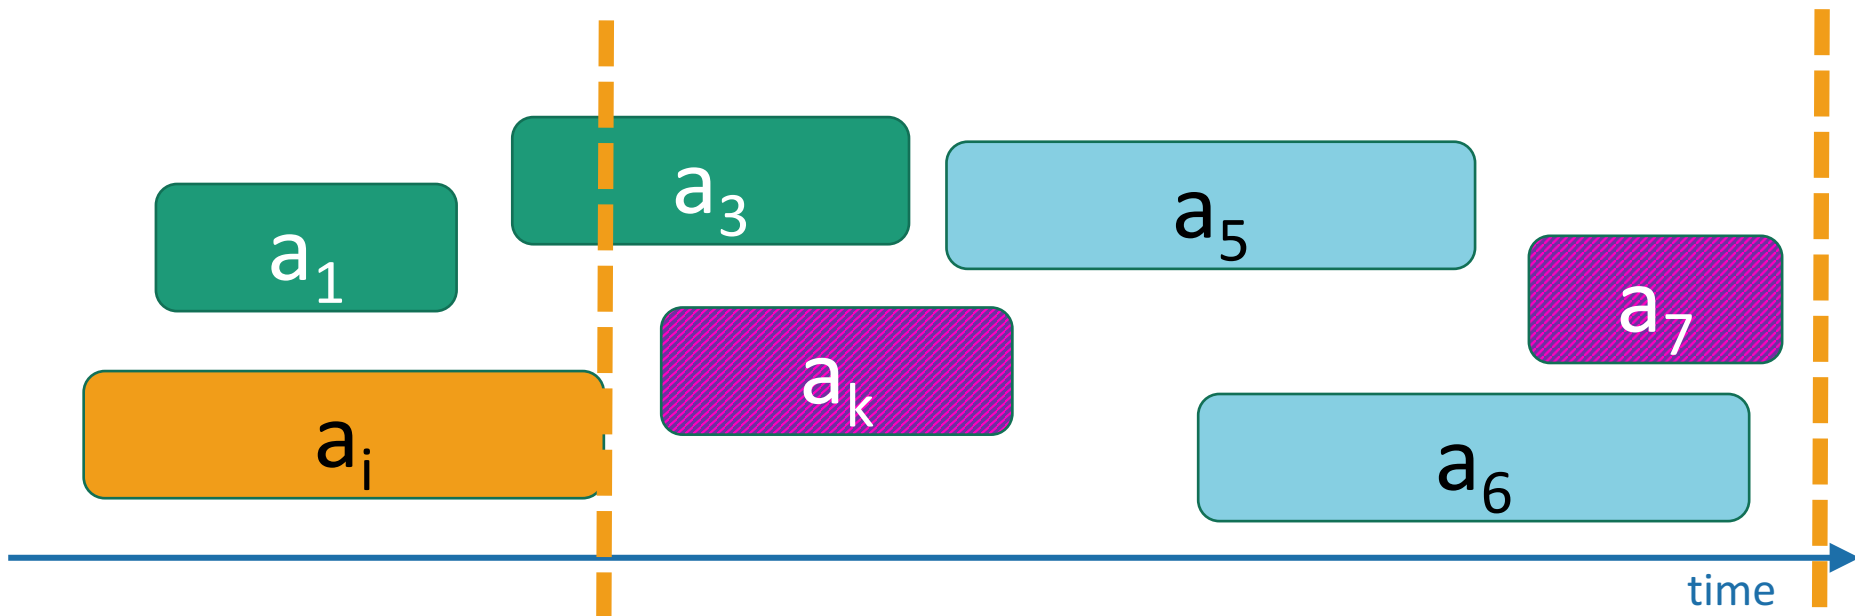
- Suppose that **this** is an optimal solution to **$A[i..n+1]$**
 - Doesn't involve a_k
- **Swap a_k in** for whatever had the smallest finishing time in **that solution**.



Proof

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .

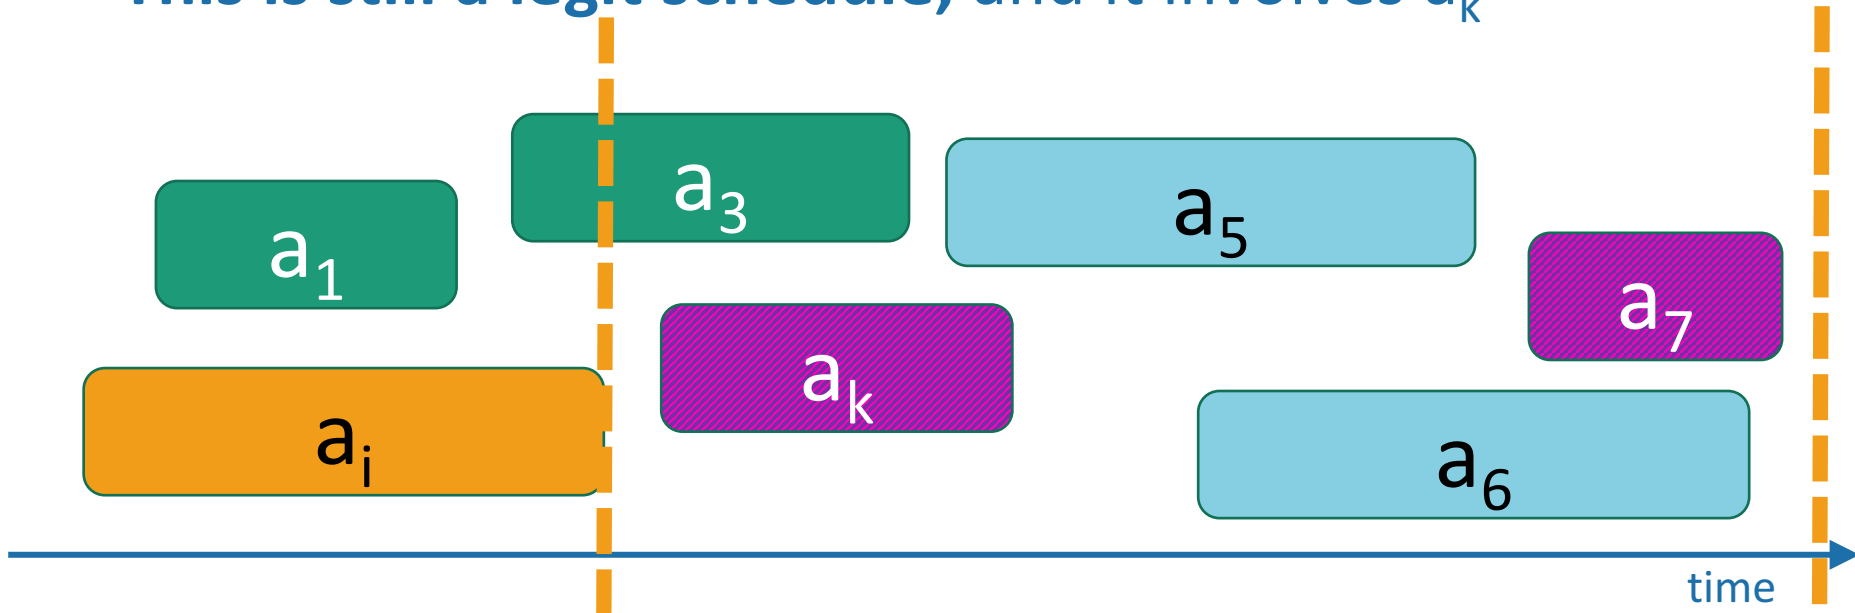
- Suppose that **this** is an optimal solution to **$A[i..n+1]$**
 - Doesn't involve a_k
- **Swap a_k in** for whatever had the smallest finishing time in **that solution**.



Proof

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .

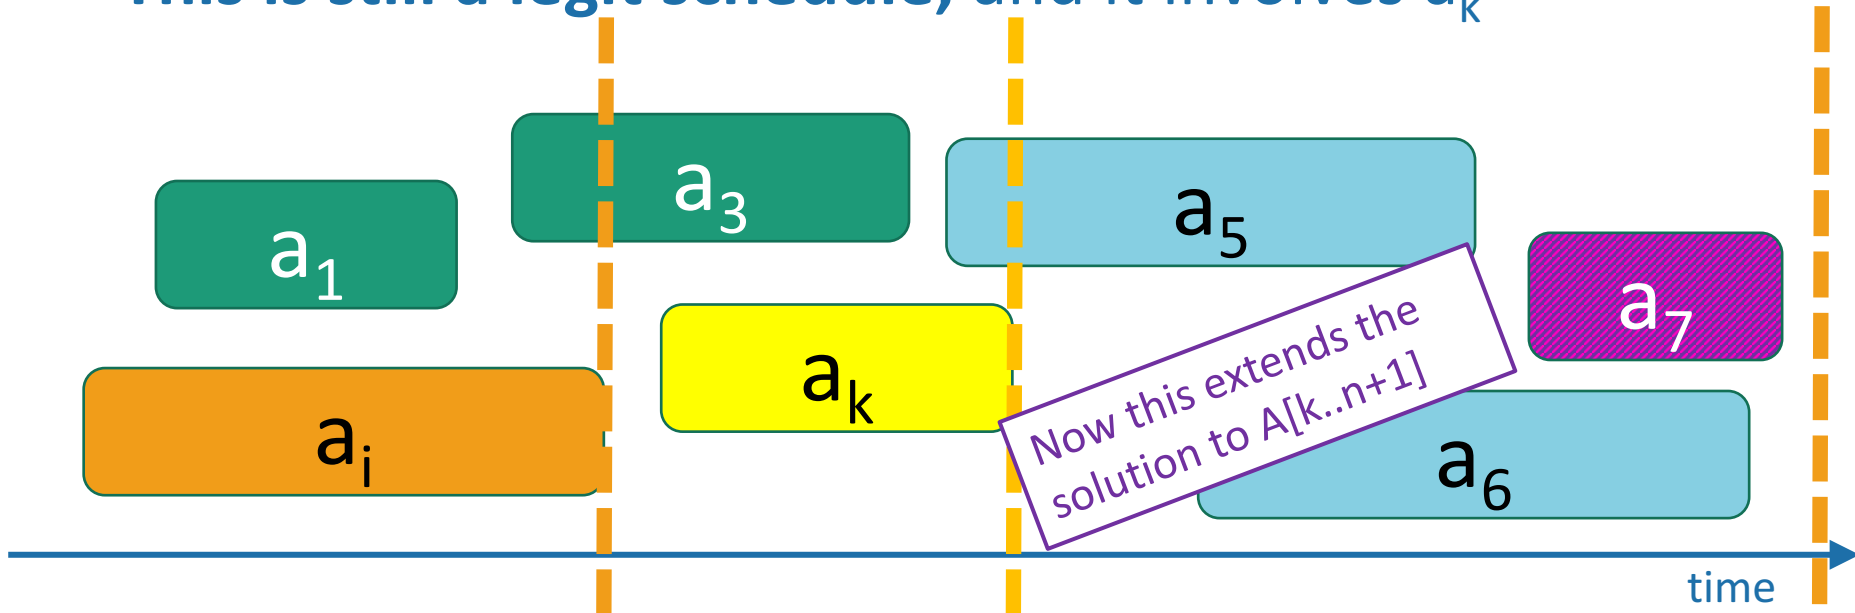
- Suppose that **this** is an optimal solution to **$A[i..n+1]$**
 - Doesn't involve a_k
- **Swap a_k in** for whatever had the smallest finishing time in **that solution**.
- **This is still a legit schedule**, and it involves a_k




Proof

- Suppose that k is the activity you can squeeze in after i **with the smallest finishing time**.
- Then there is an **optimal solution to $A[i..n+1]$** that **extends the optimal solution to $A[k..n+1]$** .


- Suppose that **this** is an optimal solution to **$A[i..n+1]$**
 - Doesn't involve a_k
- **Swap a_k in** for whatever had the smallest finishing time in **that solution**.
- **This is still a legit schedule**, and it involves a_k




This means that DP would have been wasteful.



This means that DP would have been wasteful.



This means that DP would have been wasteful.



$A[0, n+1]$ is the return value we wanted.

We should know ahead of time that it only depends on $A[2, n+1]$

This means that DP would have been wasteful.

	0					j		n+1
0								
i						A[i,j]		
n+1								

$A[0,n+1]$ is the return value we wanted.

We should know ahead of time that it only depends on $A[2,n+1]$

etc.

This means that DP would have been wasteful.

	0					j		n+1
0								
i						A[i,j]		
n+1								

$A[0,n+1]$ is the return value we wanted.

We should know ahead of time that it only depends on $A[2,n+1]$

etc.

This means that DP would have been wasteful.

	0					j		n+1
0								
i						A[i,j]		
n+1								

$A[0,n+1]$ is the return value we wanted.

We should know ahead of time that it only depends on $A[2,n+1]$

etc.

This means that DP would have been wasteful.

	0					j		n+1
0								
i						A[i,j]		
n+1								

$A[0,n+1]$ is the return value we wanted.

We should know ahead of time that it only depends on $A[2,n+1]$

etc.

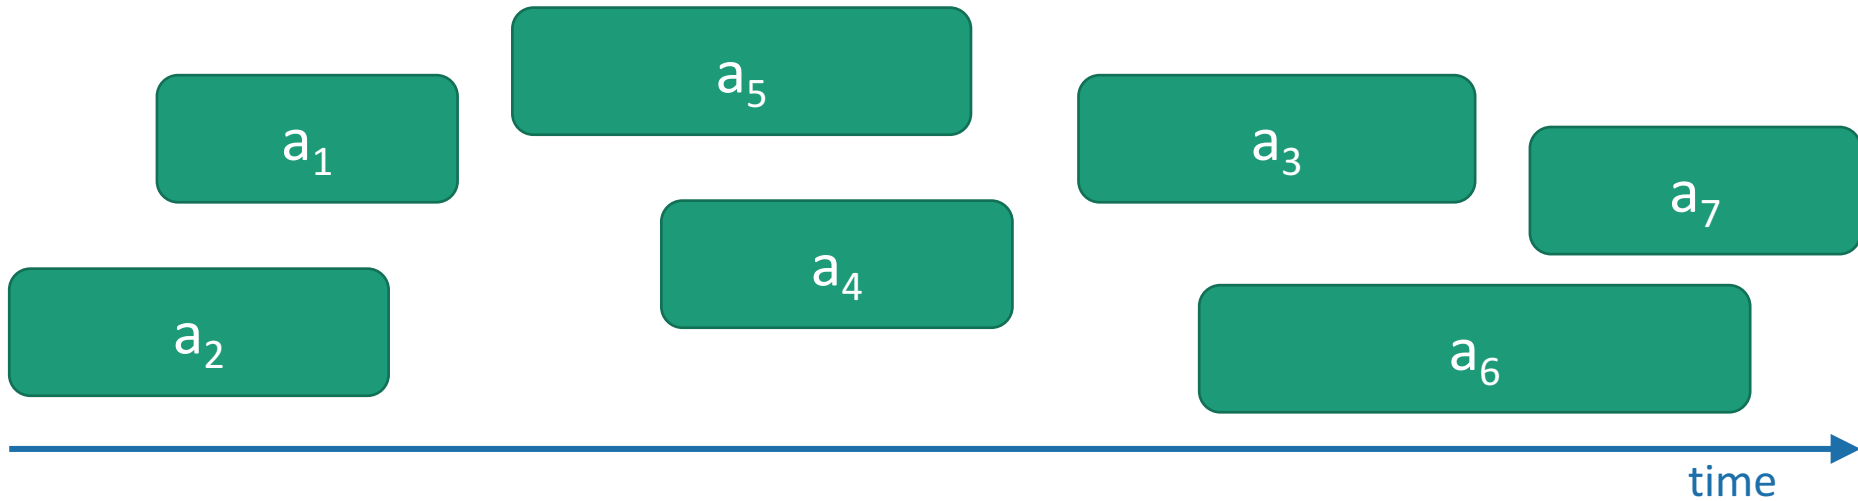
There's no reason we have to look at the whole table!

Instead, let's use this insight to make a **greedy algorithm**

- Suppose the activities are sorted by finishing time
 - if not, sort them.
- `mySchedule = []`
- **for** `k = 1,...,n`:
 - **if** I can fit in **Activity k** after the last thing in `mySchedule`:
 - `mySchedule.append(Activity k)`
- **return** `mySchedule`

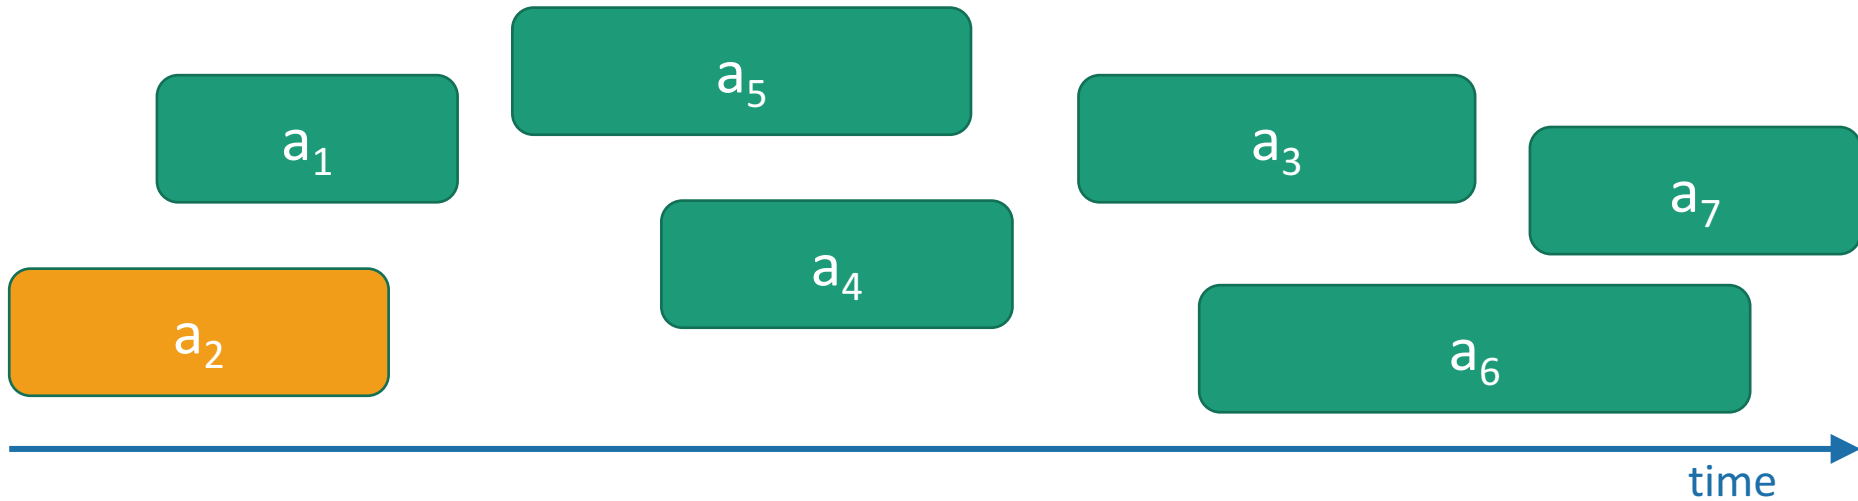
This is the same thing
we saw before

Greedy Algorithm




- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



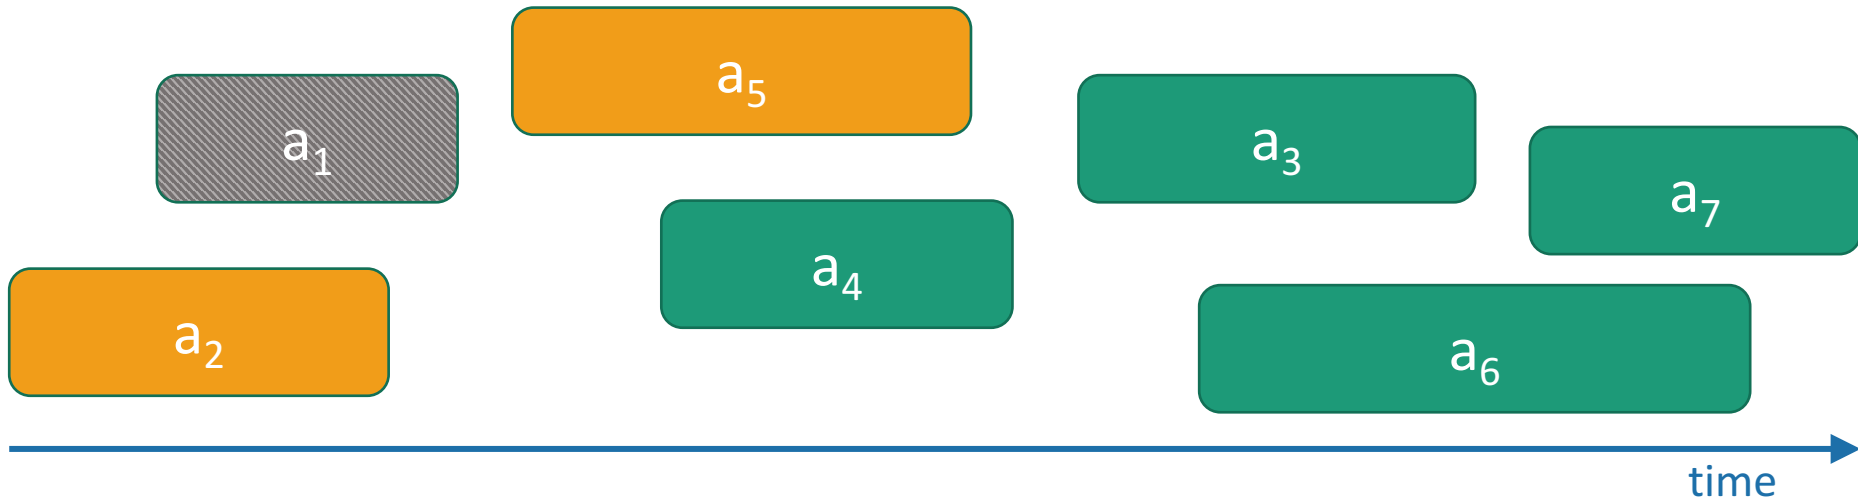
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



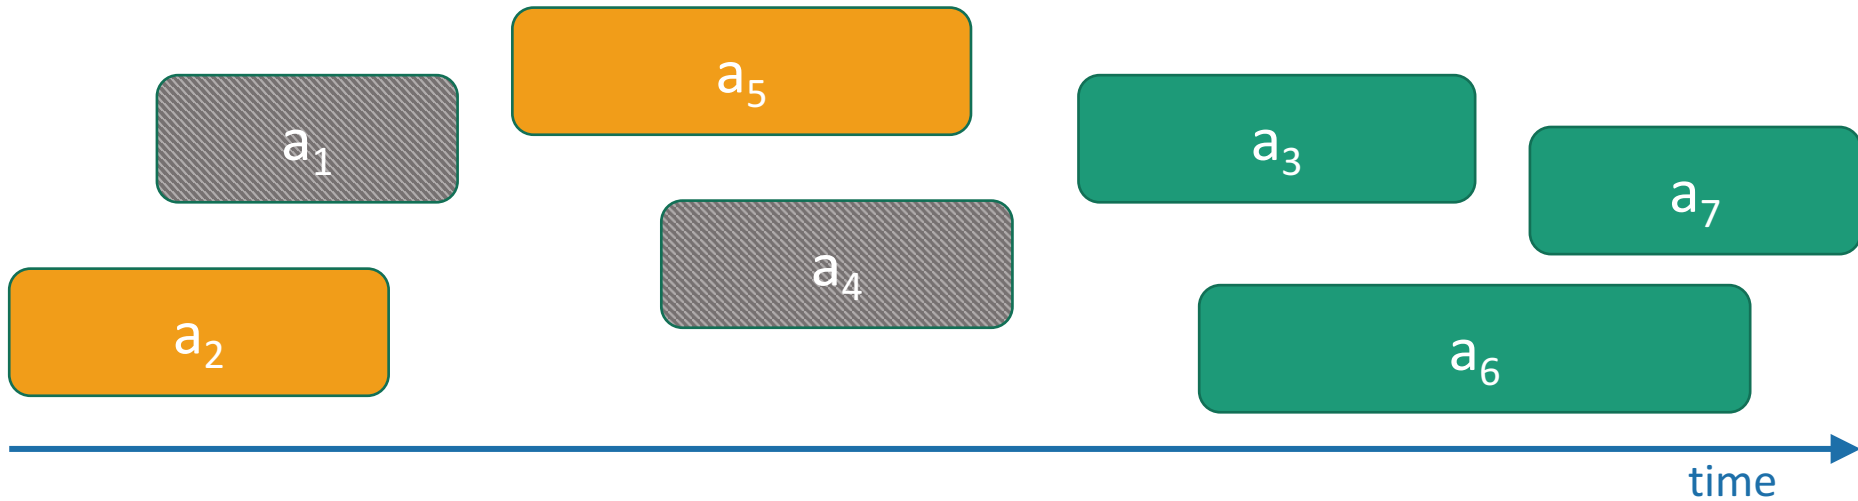
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



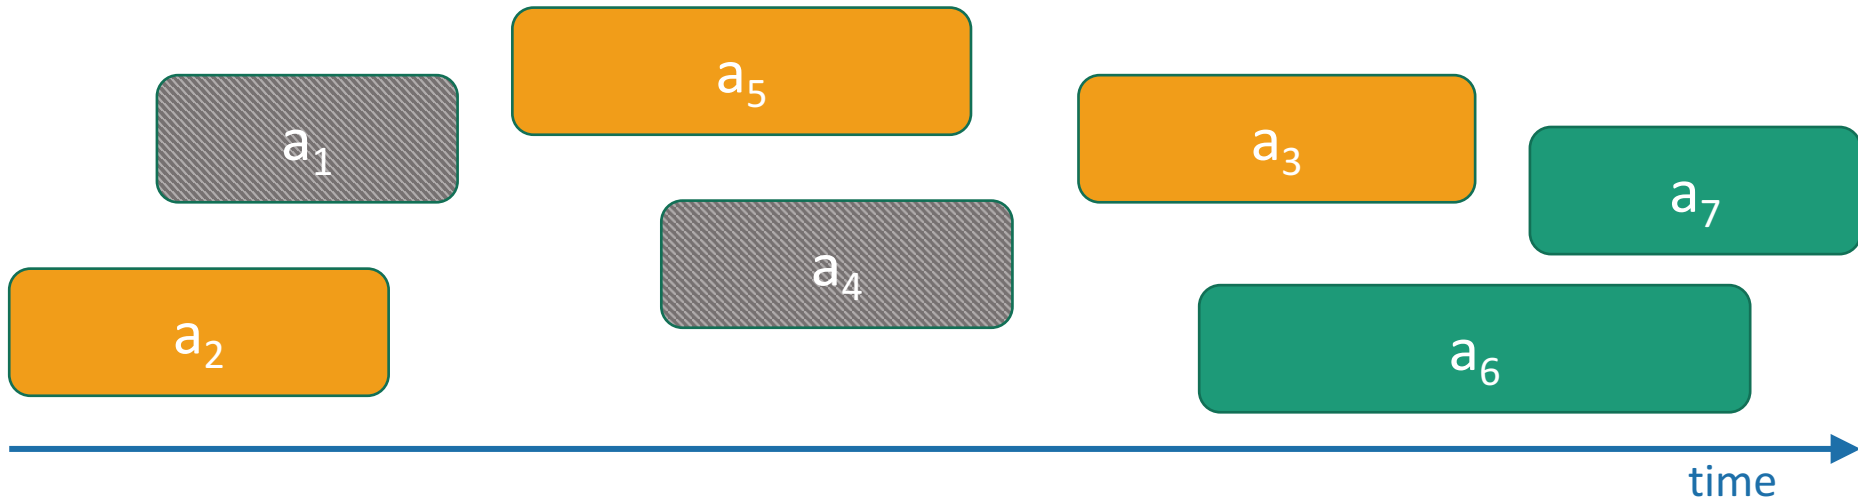
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm




- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



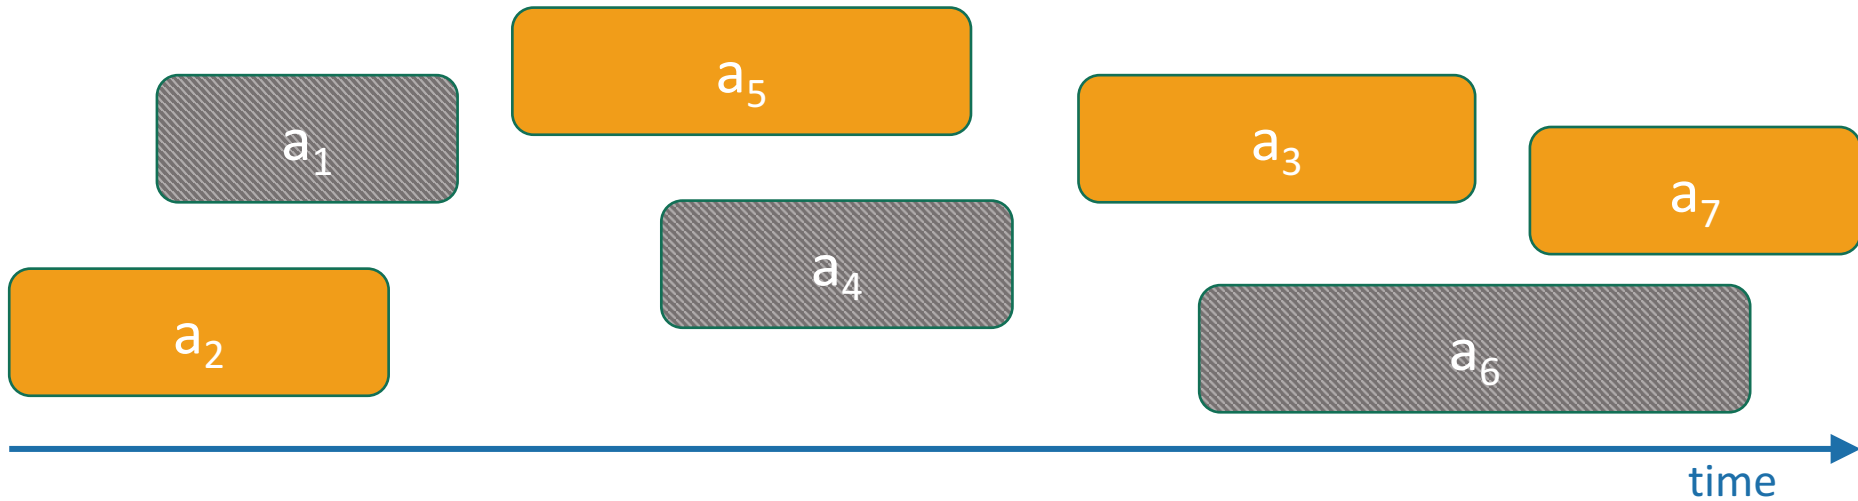
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Greedy Algorithm



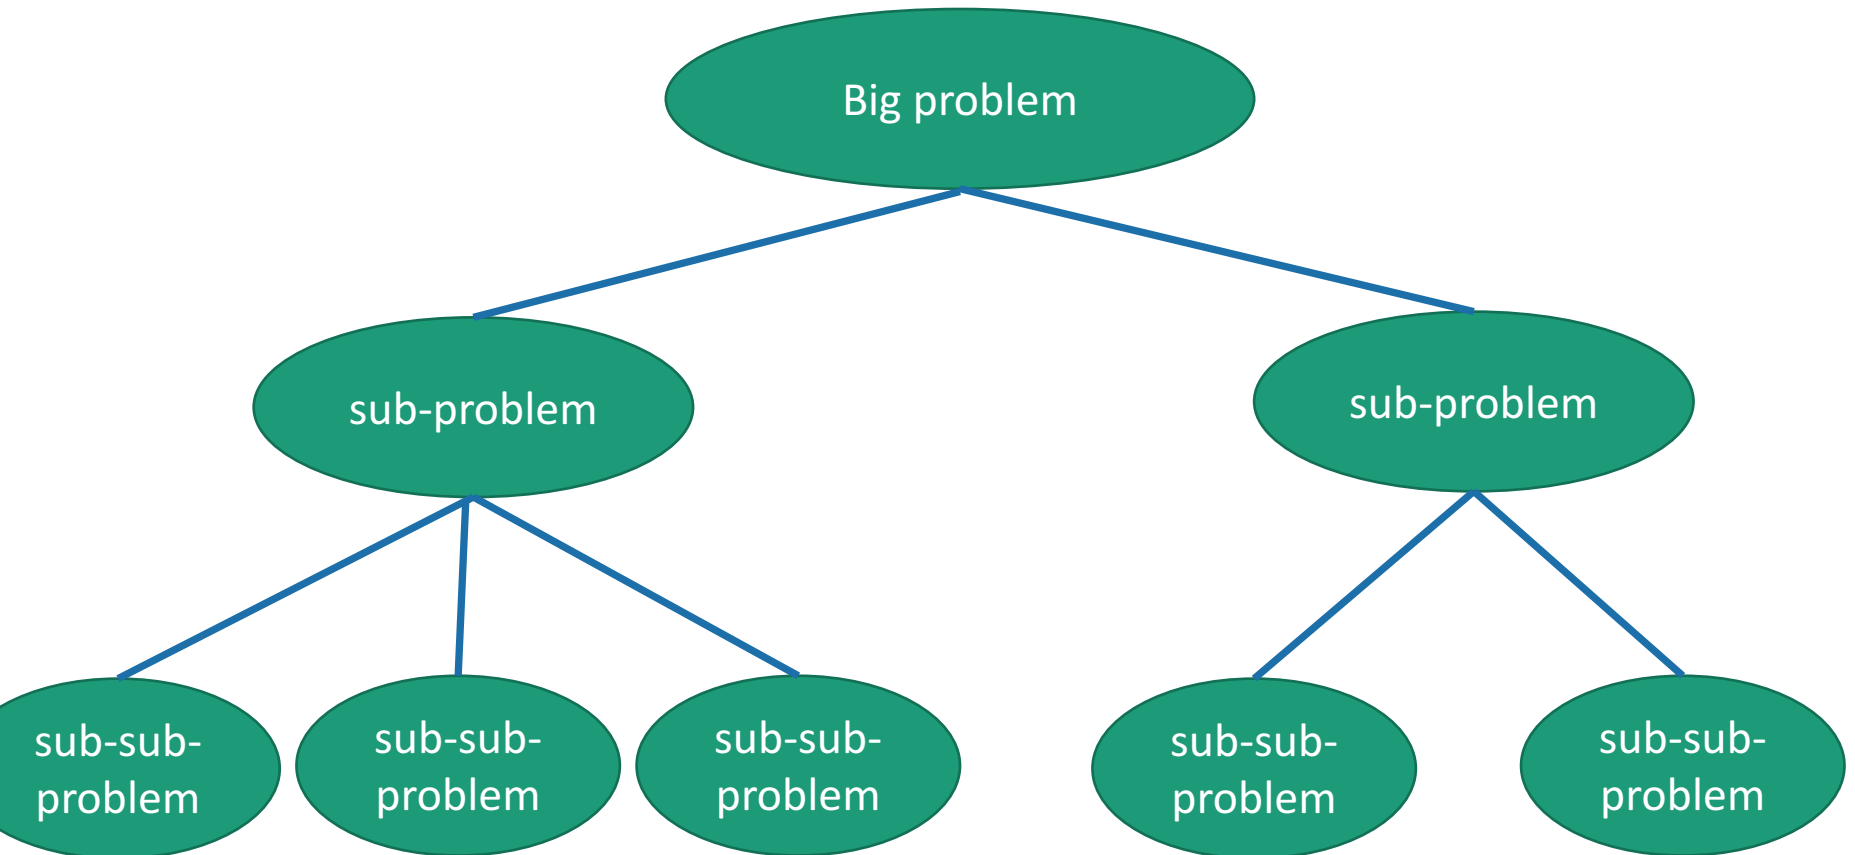
- Pick the activity you can add
 - that has the **smallest finish time**.
- Include it in your activity list.
- **Repeat.**

Why does this work?

- At each step, **we make a choice**
 - Include activity k
- We can show that this choice **will never rule out an optimal solution.**
 - Formally: There is an optimal solution to $A[i..n+1]$ that contains $A[k..n+1]$.
- So when we reach the **end of the argument:**
 - we haven't ruled out an optimal solution
 - and we only have one solution left
 - **so it must be optimal.**


Sub-problem graph view

- Divide-and-conquer:




Sub-problem graph view

- Dynamic Programming:




Sub-problem graph view

- Greedy algorithms:



Sub-problem graph view

- Greedy algorithms:



- Not only is there **optimal sub-structure**:
 - optimal solutions to a problem are made up from optimal solutions of sub-problems
- but each problem **depends on only one sub-problem**.

What have we learned?

- If we come up with a DP solution, and it turns out that we really only care about one sub-problem, then maybe we can use a greedy algorithm.
- One example was activity selection.
- In order to come up with a greedy algorithm, we:
 - Made a series of choices
 - Proved that our choices will never rule out an optimal solution.
 - Conclude that our solution at the end is optimal.

One more example

Huffman coding

- everyday english sentence

- 01100101 01110110 01100101 01110010 01111001 01100100 01100001
01111001 00100000 01100101 01101110 01100111 01101100 01101001
01110011 01101000 00100000 01110011 01100101 01101110 01110100
01100101 01101110 01100011 01100101

- qwertyui_opasdfg+hjklzxcv

- 01110001 01110111 01100101 01110010 01110100 01111001 01110101
01101001 01011111 01101111 01110000 01100001 01110011 01100100
01100110 01100111 00101011 01101000 01101010 01101011 01101100
01111010 01111000 01100011 01110110


One more example

Huffman coding

ASCII is pretty wasteful. If **e** shows up so often, we should have a more parsimonious way of representing it!


- **e**veryday **e**nglish **s**entence
- 01100101 01110110 01100101 01110010 01111001 01100100 01100001
01111001 00100000 01100101 01101110 01100111 01101100 01101001
01110011 01101000 00100000 01110011 01100101 01101110 01110100
01100101 01101110 01100011 01100101
- qwertyui_opasdfg+hjklzxcv
- 01110001 01110111 01100101 01110010 01110100 01111001 01110101
01101001 01011111 01101111 01110000 01100001 01110011 01100100
01100110 01100111 00101011 01101000 01101010 01101011 01101100
01111010 01111000 01100011 01110110

Suppose we have some distribution on characters




Suppose we have some distribution on characters

For simplicity,
let's go with this
made-up example



Try 1


- Every letter is assigned a **binary string** of one or two bits.
- The more frequent letters get the shorter strings.
- **Problem:**
 - Does 000 mean AAA or BA or AB?
some codes are prefixes of others



example text : Sleeplessness

characters	frequency	code
s	5	0
e	4	1
l	2	00
p	1	01
n	1	10

0 is the prefix of 00,01
1 is the prefix of 10




character node that has children
it's code is the prefix of another code

! all characters must be leaves in order not to be prefixes of another

Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding

- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.




10010101

Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding


- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding


- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding

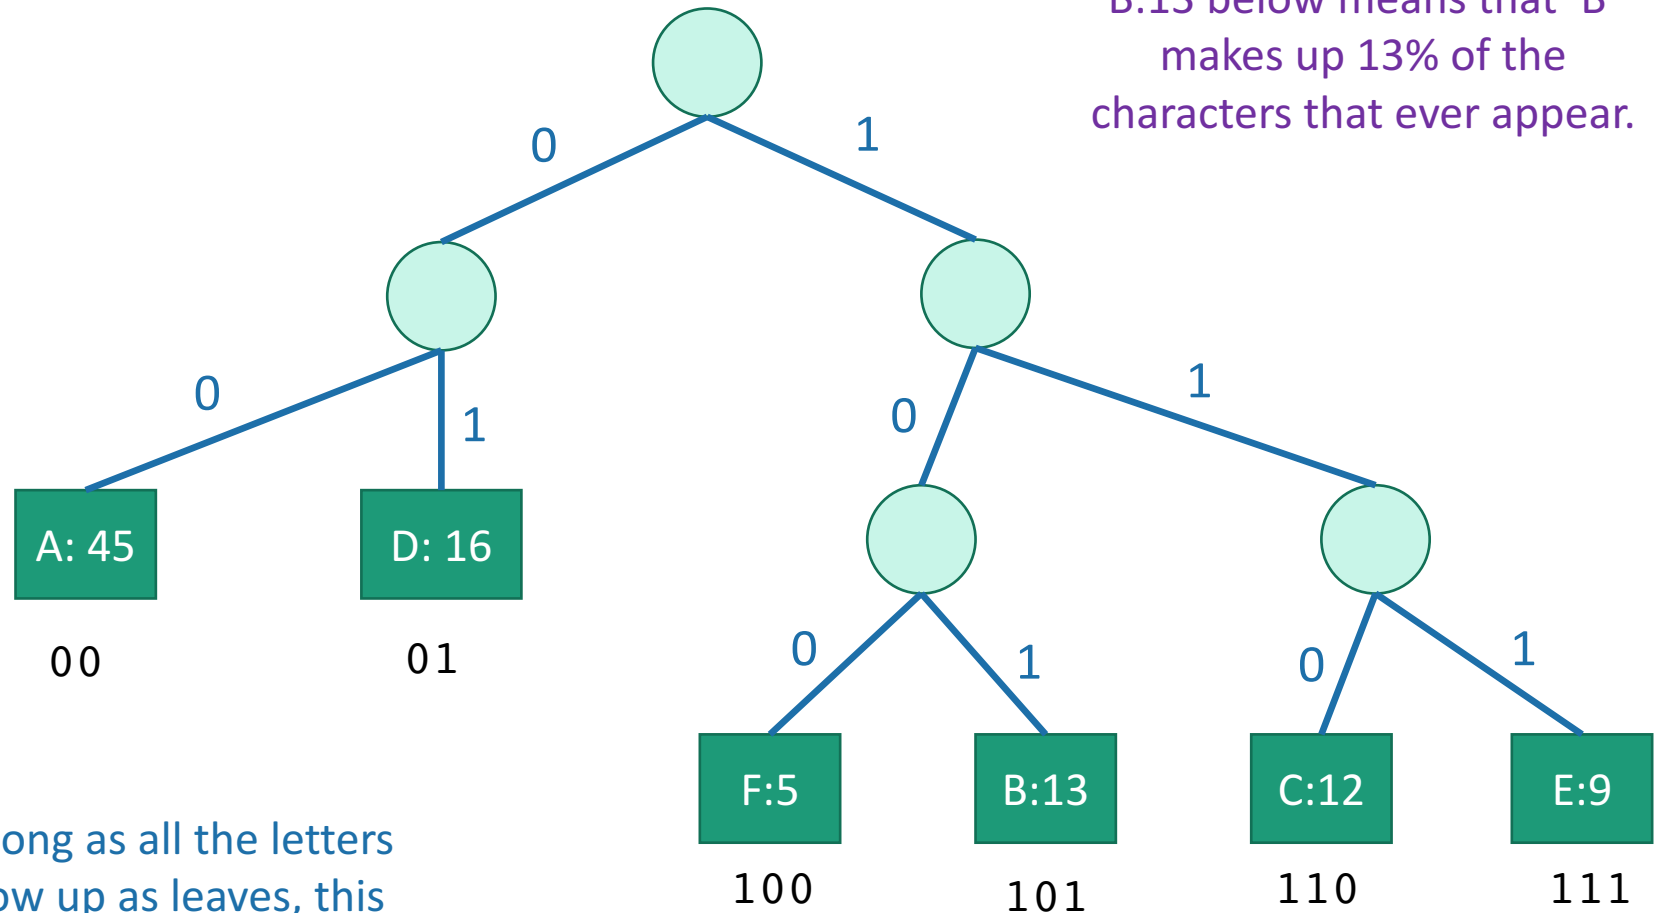
- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Question: What is the most efficient way to do prefix-free coding? (This isn't it).

A prefix-free code is a tree


B:13 below means that 'B'
makes up 13% of the
characters that ever appear.



As long as all the letters
show up as leaves, this
code is **prefix-free**.

Some trees are better than others

- Imagine choosing a letter at random from the language.
 - Not uniform, but according to our histogram!
- The **cost of a tree** is the expected length of the encoding of that letter.




Expected cost of encoding a letter with this tree:

$$2(0.45 + 0.16) + 3(0.05 + 0.13 + 0.12 + 0.09) = 2.39$$

Optimal sub-structure

- Suppose this is an optimal tree:




Then this is an optimal tree on fewer letters.

Otherwise, we could change this sub-tree and end up with a better overall tree.

In order to design a greedy algorithm

- Think about what letters belong in this sub-problem...



What's a **safe choice** to make for these lower sub-trees?

Infrequent elements!

We want them as low down as possible.


Solution

greedily build subtrees, starting with the infrequent letters




Solution

greedily build subtrees, starting with the infrequent letters




Solution

greedily build subtrees, starting with the infrequent letters



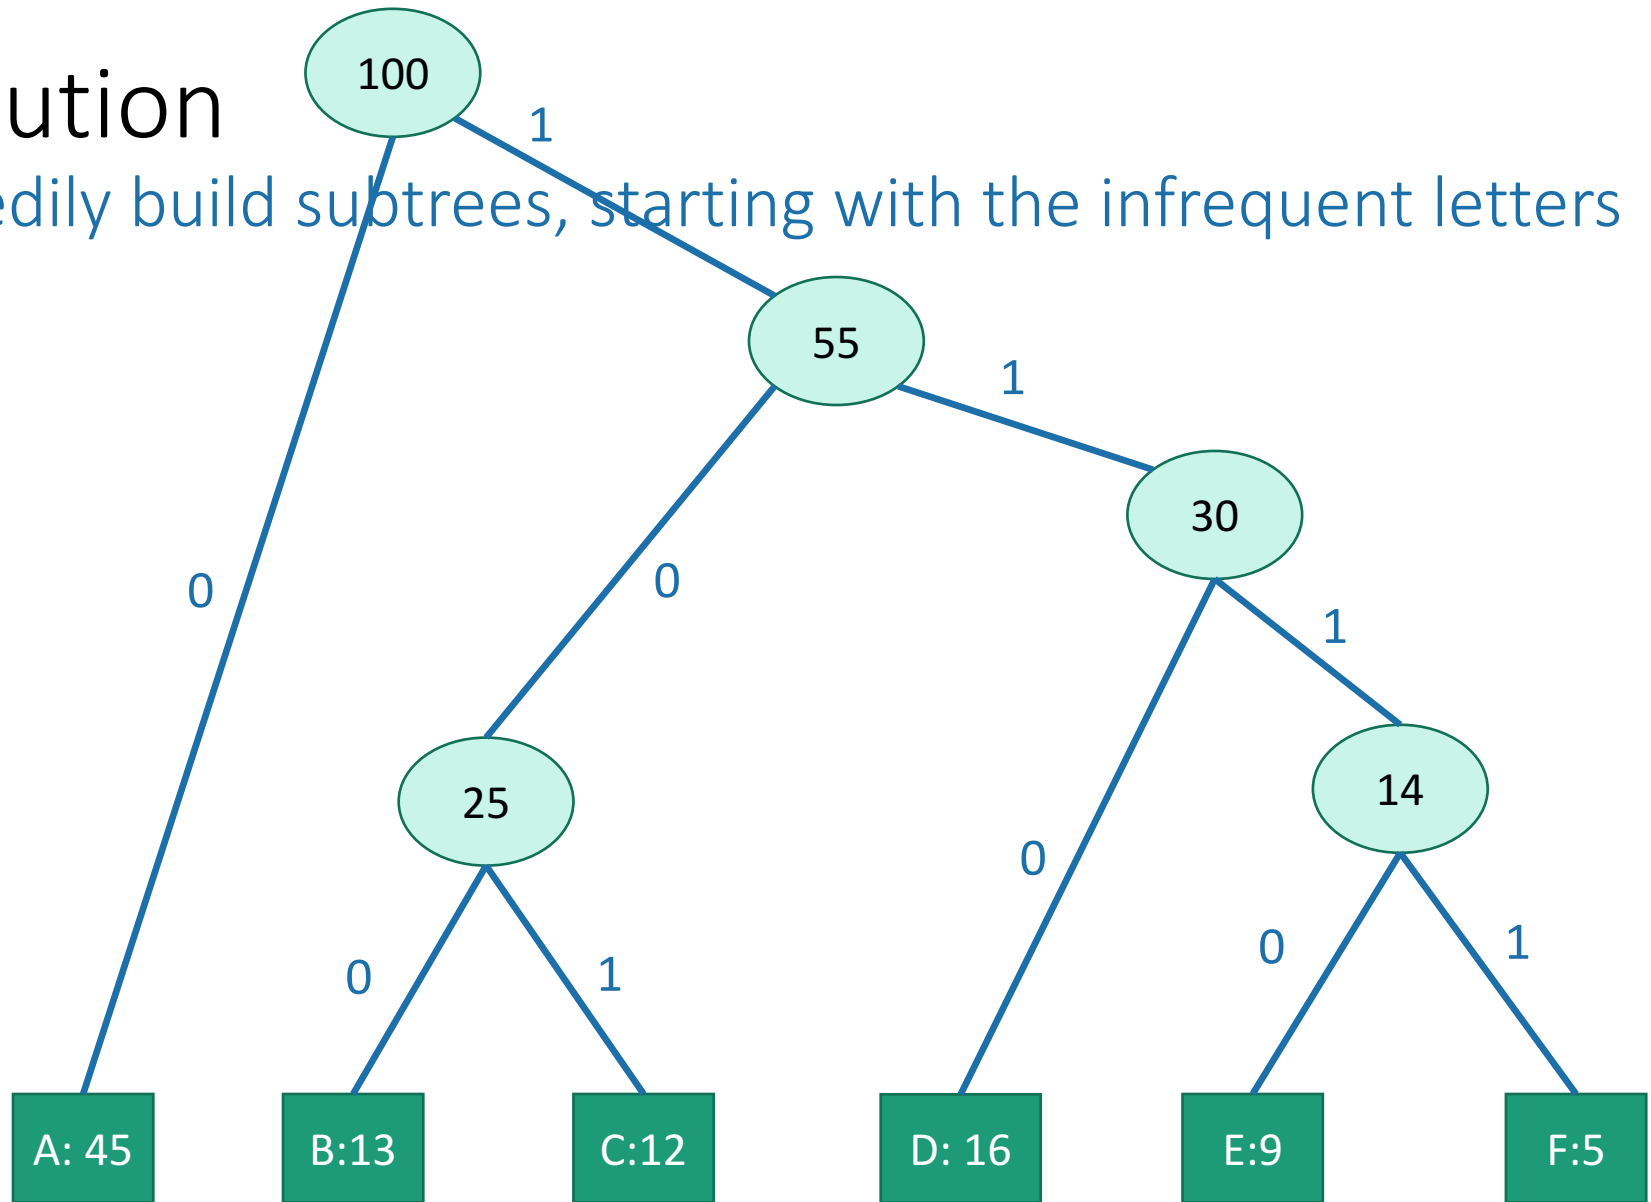
Solution

greedily build subtrees, starting with the infrequent letters



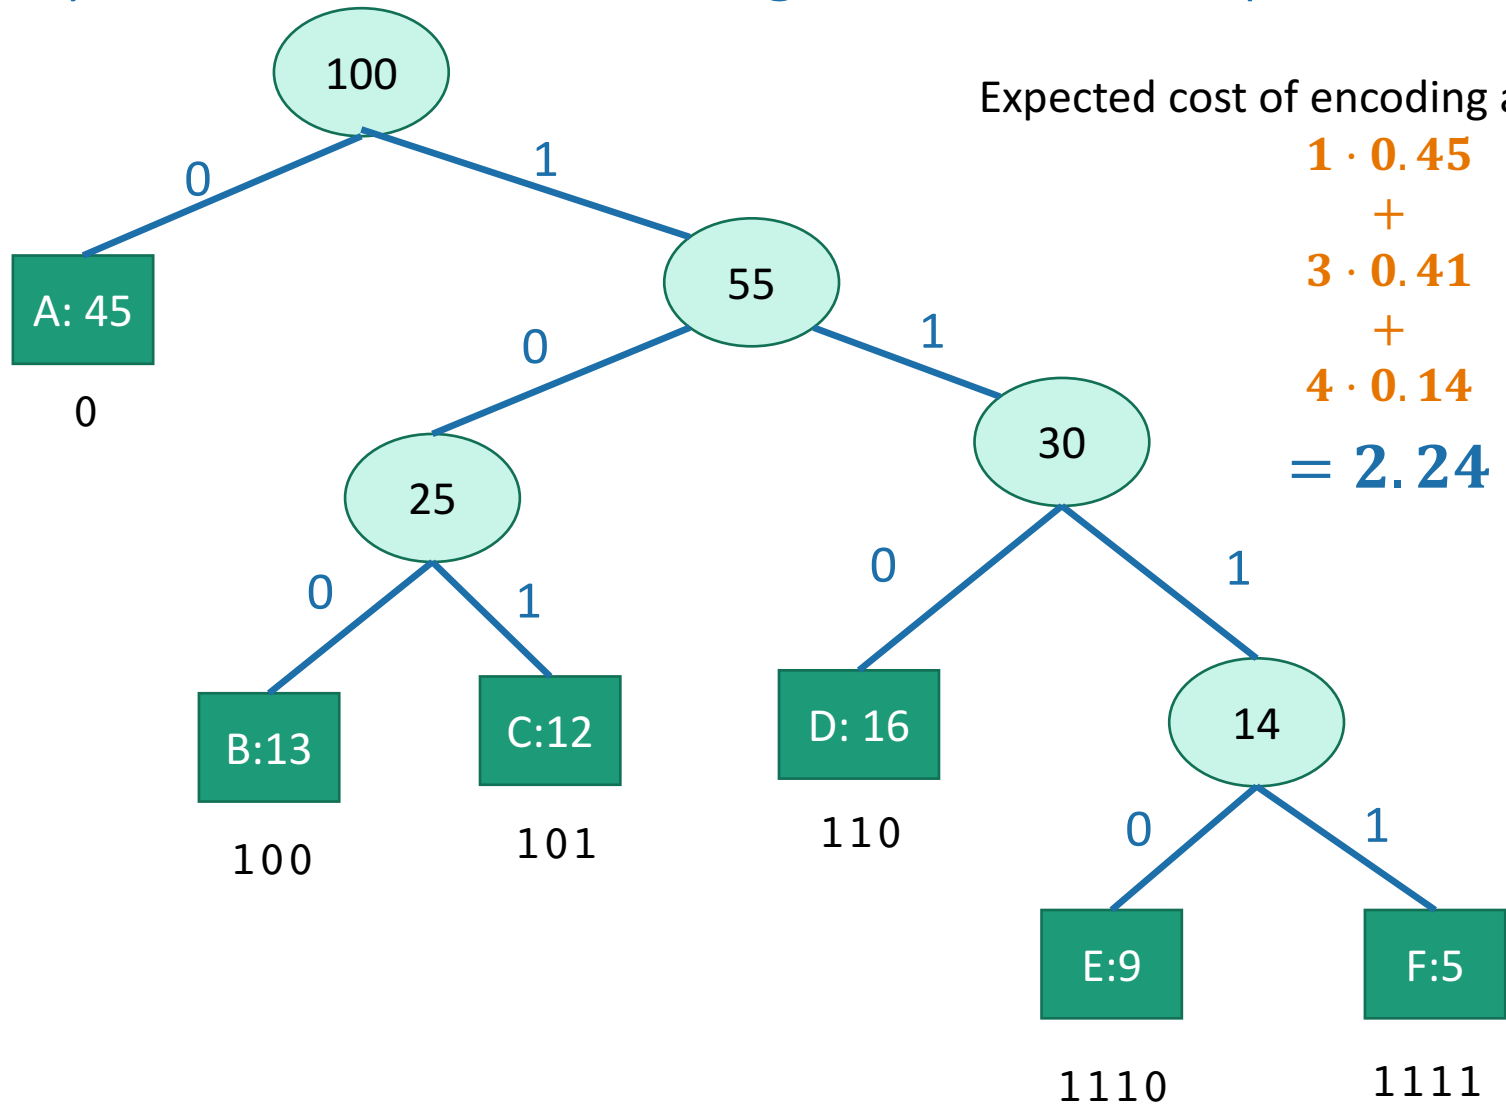
Solution

greedily build subtrees, starting with the infrequent letters



Solution

greedily build subtrees, starting with the infrequent letters



What exactly was the algorithm?


- Create a node like **D: 16** for each letter/frequency
 - The key is the frequency (16 in this case)
- Let **CURRENT** be the list of all these nodes.
- **while** len(**CURRENT**) > 1:
 - **X** and **Y** ← the nodes in **CURRENT** with the smallest keys.
 - Create a new node **Z** with **Z.key = X.key + Y.key**
 - Set **Z.left = X, Z.right = Y**
 - Add **Z** to **CURRENT** and remove **X** and **Y**
- return **CURRENT**[0]

A: 45

B: 13

C: 12

D: 16



Proof strategy

just like before

- Show that at each step, the choices we are making **won't rule out** an optimal solution.
- **Lemma:**
 - Suppose that x and y are the two least-frequent letters. Then there is an optimal tree where x and y are siblings.




Lemma

proof idea

If x and y are the two least-frequent letters, there is an optimal subtree where x and y are siblings.

- Say that an optimal tree looks like this:



Lowest-level sibling nodes: at least one of them is neither x nor y


- What happens to the cost if we swap x for a ?
 - the cost can't increase; a was more frequent than x , and we just made its encoding shorter.
- Repeat this logic until we get an optimal tree with x and y as siblings.

Lemma

proof idea

If x and y are the two least-frequent letters, there is an optimal subtree where x and y are siblings.

- Say that an optimal tree looks like this:



Lowest-level sibling nodes: at least one of them is neither x nor y

- What happens to the cost if we swap x for a ?
 - the cost can't increase; a was more frequent than x , and we just made its encoding shorter.

let's name it T and T' :


$$\text{cost}(T) - \text{cost}(T') = (p_a - p_x) \Delta_1 + (p_b - p_y) \Delta_2 \geq 0$$

Proof strategy

just like last time

- Show that at each step, the choices we are making **won't rule out** an optimal solution.
- **Lemma:**
 - Suppose that x and y are the two least-frequent letters.
Then there is an optimal tree where x and y are siblings.

Actually that's not quite enough...




Proof strategy

just like last time

Our argument before just showed that we made the right choice at the first step, when everything was a leaf. What about once we start grouping stuff?



- Show that at each step, the choices we are making **won't rule out** an optimal solution.
- Lemma:
 - Suppose that x and y are the two least-frequent letters. Then there is an optimal tree where x and y are siblings.

Actually that's not quite enough...



Lemma 2

this distinction doesn't really matter



The first thing is an optimal tree on $\{A, B, C, D, E, F\}$

if and only if

the second thing is an optimal tree on $\{A, G, H\}$

Together


- Lemma 1:
 - Suppose that x and y are the two least-frequent letters.
Then there is an optimal tree where x and y are siblings.
- Lemma 2:
 - We may as well imagine that CURRENT contains only leaves.
- These imply:
 - At each step, our choice doesn't rule out an optimal tree.

Formally, we'd use induction

After the t 'th step, we've got a bunch of current sub-trees:

- Inductive hypothesis:

- after the t 'th step,
 - there is an optimal tree containing the current subtrees as “leaves”



- Base case:

- after the 0'th step,
 - there is an optimal tree containing all the characters.


*Inductive hyp. asserts
that our subtrees can be
assembled into an
optimal tree:*

- Inductive step:


- **TO DO**

- Conclusion:

- after the last step,
 - there is an optimal tree containing this whole tree as a subtree.
- aka,
 - after the last step the tree we've constructed is optimal.

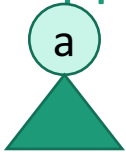



We've got a bunch of current sub-trees:




say that x and y are the two smallest.

Inductive step

- Suppose that the inductive hypothesis holds for t-1
 - After t-1 steps, there is an optimal tree containing all the current sub-trees as “leaves.”
- Want to show:
 - After t steps, there is an optimal tree containing all the current sub-trees as leaves.
- Two ingredients:
 - **Lemma 1:** If x and y are the two least-frequent letters, there is an optimal subtree where x and y are siblings.
 - **Lemma 2:** Suppose that there is an optimal tree containing  as a subtree. Then we may as well replace it with a new letter with frequency .

What have we learned?

- ASCII isn't an optimal way to encode English, since the distribution on letters isn't uniform.
- Huffman Coding is an optimal way!
- To come up with an optimal scheme for any language efficiently, we can use a **greedy algorithm**.
- To come up with a **greedy algorithm**:
 - Identify **optimal substructure**
 - Find a way to make “safe” choices that **won't rule out an optimal solution**.
 - Create subtrees out of the smallest two current subtrees.



Recap I

- Greedy algorithms!

examples:

- Activity Selection
- Huffman Coding

Recap II

- Greedy algorithms!
- Often easy to write down
 - But may be hard to come up with and hard to justify
- The natural greedy algorithm may not always be correct
- A problem is a good candidate for a greedy algorithm if:
 - it has optimal substructure
 - that optimal substructure is **REALLY NICE**
 - solutions depend on just one other sub-problem.